Energy Copilot App Product Requirements Document (PRD)

Product Name (Working)

Clarity

Daily Energy & Recovery Copilot

---

1. Problem Statement

Millions of people feel chronically tired, mentally drained, or inconsistent in energy, despite sleeping, exercising, and “doing the right things.”

Existing health and wellness apps fail because they:

- Surface metrics without meaning
- Require high effort (manual logging, complex dashboards)
- Focus on optimization, not restoring baseline well-being
- Do not explain why today feels worse than yesterday

Core Problem

This results in confusion, anxiety, and abandonment of health tools.

---

2. Product Vision

Create a daily companion app that:

- Explains why users feel the way they do
- Predicts energy dips before they happen
- Recommends small, realistic actions
- Requires less than one minute per day

The product should feel like:

---

3. Target Users

Primary ICP

- Ages 25–55
- Busy adults, knowledge workers, parents
- Generally healthy but feel “off,” tired, or inconsistent
- Overwhelmed by health data but still motivated to feel better

Secondary ICP

- Burnout-prone professionals
- Remote and hybrid workers
- Users of wearables who feel data-rich but insight-poor

Explicitly not targeting:

- Elite athletes
- Medical patients

- Biohacking or quantified-self enthusiasts

---

4. Jobs To Be Done (JTBD)

1. Understand energy
    - “Why do I feel like this today?”
2. Reduce uncertainty
    - “Is this normal or something I should worry about?”
3. Take simple action
    - “What’s the one thing I should do right now?”
4. Feel in control
    - “I can prevent this next time.”

---

5. Core Value Proposition

This product does not compete on tracking.

It competes on interpretation and clarity.

---

6. Key Features (MVP Scope)

6.1 Daily Energy Score & Explanation (Core Experience)

User sees (single primary screen):

- Energy score (1–10)
- Plain-English explanation (1–2 sentences)
- Maximum of 2–3 recommended actions

Example:

> Energy Today: 6.5 / 10
> You slept enough hours, but later than usual, and yesterday involved prolonged mental effort. This combination often leads to a mid-day energy dip.
> Try today:
> • Delay caffeine until 9:30am
> • Take a 10-minute walk before noon

No charts required to get value.

6.2 Lightweight Daily Check-Ins (≤60 seconds/day)

Purpose

Capture subjective energy and cognitive state, which is the most reliable indicator of daily well-being. These signals anchor all interpretations and explanations.

The design prioritizes speed, structure, and consistency.

Morning Check-In (Required)

Triggered via notification or app open.

Inputs:

1. “How rested do you feel right now?”
    - Scale: 1–10
2. “How motivated do you feel to start the day?”
    - Low / Medium / High

Value:

- Establishes perceived recovery baseline
- Anchors daily energy modeling
- <10 seconds to complete

Mid-Day Energy Pulse (Optional)

Triggered during predicted dip window (default: 1–3pm).

Inputs:

1. “How’s your energy right now?”
    - Low / OK / High
2. “Which best describes your state?”
    - Mentally drained
    - Physically tired
    - Distracted
    - Fine

Value:

- Validates or refines predictions
- Distinguishes mental vs physical fatigue
- Enables just-in-time guidance

Evening Reflection (Required)

Triggered in evening or pre-bed.

Inputs (single-tap):

1. “What drained you most today?”
    - Poor sleep
    - Work / mental load
    - Physical exertion
    - Emotional stress
    - Poor meals or timing
    - Unknown
2. “How did today compare to expectations?”
    - Better / Same / Worse

Value:

- Captures perceived causality
- Trains explanation accuracy
- Improves future forecasts

Data Handling

- No free text required
- Structured, categorical inputs only

- Skippable without penalty
- Stored as time-series signals tied to daily outcomes

6.3 Passive & Integrated Data Sources (Opt-In)

Purpose

Improve explanation quality using low-friction passive signals, while ensuring the product delivers value even without integrations.

All integrations are optional, permission-based, and transparent.

6.3.1 Sleep Data

Primary Integrations:

- Apple Health (iOS)
- Google Health Connect (Android)

Data used:

- Total sleep duration
- Bedtime and wake time
- Sleep consistency (variance over days)

Fallback (No Wearable):

- Phone lock/unlock patterns
- Alarm usage
- Screen activity near bedtime

Rationale:

Sleep timing consistency is more predictive of energy than sleep stages.

6.3.2 Calendar & Cognitive Load

Integrations:

- Google Calendar
- Apple Calendar
- Microsoft Outlook (Phase 2)

Data accessed (explicitly limited):

- Meeting start/end times
- Meeting density
- Gaps between meetings
- Late-day meetings

Explicitly not accessed:

- Titles
- Attendees
- Notes
- Content

Rationale:

Meeting density and lack of recovery gaps are major hidden energy drains.

6.3.3 Movement & Sedentary Behavior

Integrations:

- Apple Health
- Google Health Connect
- Device motion sensors (fallback)

Data used:

- Step count ranges (low / medium / high)
- Long sedentary blocks
- Presence or absence of movement

Rationale:

Explains afternoon crashes and enables low-effort recovery actions.

6.3.4 Habit & Context Flags (Binary)

Collected via 1-tap prompts.

Signals:

- Late caffeine (after user-defined cutoff): Yes / No
- Skipped or delayed meals: Yes / No
- Alcohol previous evening: Yes / No

Collection points:

- Evening reflection

- Occasional reminders if missing

Rationale:

Binary inputs reduce friction while capturing sufficient signal.

6.3.5 Optional Wearable Enhancements (Phase 2)

Supported:

- Apple Watch
- Oura
- Fitbit
- Whoop

Used for:

- Sleep consistency validation
- Recovery trend confirmation

Not surfaced to users:

- HRV
- Raw biometrics
- Readiness scores

Wearables enhance confidence, not dependency.

6.3.6 Data Fusion & Explanation Engine (High-Level)

The system builds a personal baseline within 5–7 days.

It evaluates:

- Short-term fatigue (24–48h)
- Accumulated recovery debt
- Cognitive load patterns
- Disruptive habits

Outputs:

- Daily energy score
- Plain-English explanation
- 2–3 context-aware actions

Uncertainty is explicitly communicated to maintain trust.

6.4 Weekly Energy Pattern Insights

Delivered once per week.

User sees:

- Top 2 energy drains
- Top 2 energy supports
- One behavior to experiment with next week

Example:

6.5 Energy Forecasting (Advanced Feature)

- Predicts high-fatigue days
- Warns users in advance
- Suggests preventative actions

Example:

6.6 Feedback Loop & Model Calibration

After explanations:

- “Did this match how you felt?”
    - Yes / Not really

This feedback:

- Improves accuracy
- Builds transparency
- Keeps users in the loop

---

7. Out of Scope (MVP)

- Medical diagnostics or claims

- Calorie or macronutrient tracking
- Detailed biometric dashboards
- Social or community features
- Therapy, coaching, or clinical interventions

---

8. User Experience Principles

1. Explain, don’t overwhelm
2. One screen answers the question
3. Actions must be achievable today
4. Admit uncertainty
5. Privacy-first by design

---

9. Success Metrics

Engagement

- DAU / WAU
- Weekly retention ≥ 35–40%
- Avg daily usage < 90 seconds

Value Validation

- % of users reporting improved understanding of energy

- % following at least one recommendation per day
- Self-reported energy improvement over 14 days

---

10. Privacy, Trust & Compliance

- No medical or diagnostic language
- Clear consent for every integration
- Explicit explanation of data usage
- Easy opt-out for any data source
- Compliance with platform health data policies

—----------

This is a **serious, well-thought-through mandate**. I'll break it down the way a founder / product-led client would be thinking, not as a feature list — so you can align execution, tradeoffs, and conversations.

I'll cover:

1. What problem the client is *really* solving
2. Why this product exists *now* (timing rationale)
3. The hidden strategy behind the PRD
4. Customer use cases that matter (day-to-day reality)
5. Product viability (what makes this work or fail)
6. What the client will judge you on (success ≠ features)

7. How this should shape your execution decisions

---

1. The Real Problem (Beyond “Low Energy”)

This app is **not about energy**.

It’s about **cognitive uncertainty**.

Most users are asking:

- *“Is something wrong with me?”*
- *“Why does yesterday feel easier than today?”*
- *“Am I burning out or just having a bad day?”*

Key Insight

Energy fluctuation is normal — **uncertainty about it is not**.

The client is trying to eliminate:

- Anxiety caused by unexplained fatigue
- Self-blame (“I’m lazy”)
- Tool fatigue from over-tracking

This is why:

- Subjective inputs are prioritized over biometrics
- Explanations matter more than scores
- Charts are explicitly discouraged

👉 The product is a **sense-making tool**, not a health tracker.

---

2. Why This Product Exists *Now*

Market Timing Signals the Client Is Responding To

1. **Wearable Saturation**
    - Users already have data
    - They don't know what it means
    - "Readiness" scores feel arbitrary and opaque
2. **Burnout Normalization**
    - Being tired is now "normal"
    - People don't want another optimization tool
    - They want reassurance and guidance
3. **LLMs Make Interpretation Cheap**
    - Explanation used to require coaches or experts
    - Now it can be automated — if done carefully
    - The PRD is designed to *avoid hallucination risk*
4. **Low Trust in Health Claims**
    - Explicit avoidance of medical language
    - Heavy emphasis on uncertainty communication

👉 The client believes **interpretation is now the new product moat**.

---

3. The Hidden Strategy Behind the PRD

This PRD is deceptively simple. Underneath, the client is making **very deliberate bets**.

Bet #1: Subjective Data Is the Truth Layer

The most important signals:

- "How rested do you feel?"

- “How drained did today feel?”

These:

- Anchor model outputs
- Prevent over-fitting to passive data
- Keep users feeling *seen*

This is why:

- No free text (to avoid cognitive load)
- No over-collection
- Check-ins are ritualized

👉 The app trains *with* the user, not *on* the user.

Bet #2: Explanations Build Retention, Not Scores

Most apps lead with:

> “Here’s your score”

This app leads with:

> “Here’s why today felt the way it did”

Scores are **secondary**. The explanation is the product.

That’s why:

- 1–2 sentences max
- Plain English only
- Explicit uncertainty allowed

👉 If the explanation feels right, the user stays.
Even if the score is wrong.

Bet #3: Small Wins Beat Optimization

Actions are:

- Deliberately boring
- Immediately achievable
- Non-heroic

The client is avoiding:

- “Do 10,000 steps”
- “Fix your sleep hygiene”
- “Optimize HRV”

Instead:

- Delay caffeine
- Walk for 10 minutes
- Eat earlier

👉 This is **behavioral compliance design**, not health advice.

---

4. Core Customer Use Cases (Real Life)

These are the *actual* moments the app is built for.

Use Case 1: “Why Am I Dragging Today?”

Morning:

- User feels off
- Opens app
- Sees explanation that matches intuition

Outcome:

- Relief, not motivation
- Reduced anxiety
- Trust established

Use Case 2: “Prevent the Afternoon Crash”

Mid-day:

- App predicts dip
- Nudges user before it happens

Outcome:

- User feels supported, not nagged
- Perceives app as proactive
- Starts checking *before* feeling bad

Use Case 3: “Was Today Normal?”

Evening:

- Reflection helps externalize blame
- Causal framing (“mental load”, “late meals”)

Outcome:

- Emotional regulation
- Improved sleep
- Higher next-day engagement

Use Case 4: “I Want Patterns, Not Advice”

Weekly insight:

- Only top 2 drains + supports
- No trend dashboards

Outcome:

- User experiments voluntarily
- Feels agency without pressure

---

5. Product Viability Analysis

Why This Can Work

✅ **Low daily friction**

- <60 seconds/day is realistic
- No guilt for skipping

✅ **Trust-first design**

- Admits uncertainty
- Avoids medical claims
- Explains data usage clearly

✅ **Strong emotional payoff**

- Feeling understood beats quantified feedback
- Rare in health apps

Where It Can Fail (Execution Risks)

❌ Over-explanation

- If explanations become verbose or generic
- If LLM output feels templated

❌ Feature creep

- Adding charts, trends, biometrics
- Diluting the “one screen clarity” promise

❌ False precision

- Overconfident predictions
- Implied medical conclusions

The PRD is designed to **protect against these failures** — your execution must respect that.

---

6. What the Client Will *Actually* Evaluate You On

Not:

- Number of integrations
- Model sophistication
- UI polish

But:

1. **Does the explanation feel human?**
2. **Does the app feel calming, not demanding?**
3. **Can a user get value with zero integrations?**
4. **Does it reduce cognitive load, not add to it?**

If you nail those, the rest is secondary.

---

7. How This Should Shape Your Execution

As an Executor, Your North Stars Are:

- **Clarity > Accuracy**
- **Trust > Intelligence**
- **Relief > Motivation**
- **Interpretation > Tracking**

Practical Implications for Build Decisions

- Start with **hard-coded explanation templates** + light LLM personalization
- Delay full forecasting until explanation accuracy feels right
- Build feedback loops early (“Did this match?”)
- Instrument *emotional* success, not just engagement

Final Framing (Use This Internally)

> This app is not trying to make users better humans.
> It’s trying to make bad days feel understandable and manageable.

If you align execution to *that belief*, you’ll be aligned with the client’s real goal.